

Zhen Peng

Email: hi.pengzhen@gmail.com

Phone: +1 (757) 378-8642

EDUCATION

Ph.D. in Computer Science

08/2016 – 01/2023

Department of Computer Science, College of William & Mary, Williamsburg, VA

Advisor: Prof. Bin Ren

M.S. in Computer Software and Theory

09/2013 – 06/2016

Department of Computer Science, Huaqiao University, Xiamen, China

Advisor: Prof. Tian Wang

B.E. in Computer Science and Technology

09/2009 – 06/2013

Department of Computer Science, Huaqiao University, Xiamen, China

RESEARCH EXPERIENCE

Post Doctorate

04/2023 – Present

Pacific Northwest National Laboratory (PNNL), Richland, WA

Automatic Code Generation for Graph Algorithms in Linear Algebra Expressions

- Design and extend intermediate representation to support parallel graph kernel in compiler COMET.
- Speedup the generated sparse kernel, such as SpGEMM

Research Assistant

08/2017 – 06/2022

Department of Computer Science, College of William & Mary, VA

Accelerate Deep Neural Network Inference on Edge Devices

- Analyze inference procedure of TensorFlow Lite for Micro on microcontroller units (MCU).
- Speed up the inference procedure by tuned loop unrolling and customized quantization method.

Efficient Parallelization of Graph-based Approximate Nearest Neighbors Search (ANNS)

- Analyze and parallelize the best-first search algorithm for ANNS on the graph-based index.
- Reduce the intra-query latency on CPUs by a tailored parallelism scheme and synchronization mechanism.

Parallelizing Pruned Landmark Labeling: Dealing with Dependencies in Graph Algorithms

- Analyze and parallelize the sequential 2-hop labeling for shortest distance queries in large graphs.
- Reduce the query latency on CPUs using the parallel algorithm that breaks the dependency.

Efficient Parallelization of Graph Processing on Emerging Many-core Architectures

- Design the graph processing system for typical graph algorithms such as BFS to tap into many-core CPUs.
- Achieve good performance and scalability by be aware of data locality, load balance, and update conflicts.

INTERNSHIP EXPERIENCE

PhD Intern

06/2022 – 04/2023

Pacific Northwest National Laboratory (PNNL), Richland, WA

Redundancy-Aware Code Generation for Sparse Tensor Expressions

- Detect the redundancy in code generation for sparse tensor expressions.
- Implement partial fusion algorithm in the MLIR-based compiler COMET.

Software Engineering Intern

04/2021 – 09/2021

Kuaishou, US R&D Center, Palo Alto, CA

Automate the Model Implementation to TensorRT

- Translate models from TVM Relay IR to TensorRT Python code.
- Speed up the model deployment procedure and reduce the labor costs.

Accelerate Inference Through Operation Fusion in Convolutional Neural Network (CNN)

- Try to add customized operation fusion pass in TVM.

- Transform fused computational graph to TensorRT C++ code to accelerate the inference.

PATENT

- [1] “Multi-Level Intermediate Representation Decoder for Heterogeneous Platforms,” U.S. Patent Application No. 17/524,619, Filing date: November 11, 2021

PUBLICATIONS

- [1] **Zhen Peng**, Rizwan A. Ashraf, Luanzheng Guo, Ruiqin Tian, and Gokcen Kestor, “Automatic Code Generation for High-Performance Graph Algorithms,” *The 32nd International Conference on Parallel Architectures and Compilation Techniques (PACT 2023)*, October 21-25, 2023, Vienna, Austria.
- [2] **Zhen Peng**, Minjia Zhang, Kai Li, Ruoming Jin, and Bin Ren, “iQAN: Fast and Accurate Vector Search with Efficient Intra-Query Parallelism on Multi-Core Architectures,” *The 28th ACM SIGPLAN Annual Symposium on Principles and Practice of Parallel Programming (PPoPP 2023)*, February 25-March 1, 2023, Montreal, Canada.
- [3] **Zhen Peng**, Minjia Zhang, Kai Li, Ruoming Jin, and Bin Ren, “Speed-ANN: Low-Latency and High-Accuracy Nearest Neighbor Search via Intra-Query Parallelism,” *arXiv:2201.13007*, 2022.
- [4] Qihan Wang, **Zhen Peng**, Bin Ren, Jie Chen, and Robert G. Edwards, “MemHC: An Optimized GPU Memory Management Framework for Accelerating Many-body Correlation,” *ACM Transactions on Architecture and Code Optimization (TACO)*, Volume 19, Issue 2, No. 24, pp 1-26, June 2022.
- [5] Ruoming Jin*, **Zhen Peng***, Wendell Wu, Feodor Dragan, Gagan Agrawal, and Bin Ren, “Parallelizing Pruned Landmark Labeling: Dealing with Dependencies in Graph Algorithms,” *The 34th ACM International Conference on Supercomputing (ICS 2020)*, June 29-July 2, 2020, Online. (* Equal contribution)
- [6] Yu Chen, Ivy Peng, **Zhen Peng**, Xu Liu, and Bin Ren, “ATMem: Adaptive Data Placement in Graph Applications on Heterogeneous Memories,” *International Symposium on Code Generation and Optimization (CGO 2020)*, February 22-26, 2020, San Diego, CA, USA.
- [7] Ruoming Jin, **Zhen Peng**, Wendell Wu, Feodor Dragan, Gagan Agrawal, and Bin Ren, “Pruned Landmark Labeling Meets Vertex Centric Computation: A Surprisingly Happy Marriage!” *arXiv:1906.12018*, 2019.
- [8] **Zhen Peng**, Alexander Powell, Bo Wu, Tekin Bicer, and Bin Ren, “GraphPhi: Efficient Parallel Graph Processing on Emerging Throughput-oriented Architectures,” *International conference on Parallel Architectures and Compilation Techniques (PACT 2018)*, November 1-4, 2018, Limassol, Cyprus.

PUBLICATIONS BEFORE PH.D.

- [1] Tian Wang, **Zhen Peng**, Sheng Wen, Weijia Jia, Yiqiao Cai, Hui Tian, and Yonghong Chen. “Reliable Wireless Connections for Fast-Moving Rail Users Based on a Chained Fog Structure.” *Information Sciences (Inf. Sci.)*, 379: 160-176, 2017.
- [2] Tian Wang, **Zhen Peng**, Chen Wang, Yiqiao Cai, Yonghong Chen, Hui Tian, Junbin Liang, and Bineng Zhong. “Extracting Target Detection Knowledge Based on Spatio-temporal Information in Wireless Sensor Networks.” *International Journal of Distributed Sensor Networks (IJDSN)*, 2016 (doi:10.1155/2016/5831471), 2016.
- [3] Tian Wang, **Zhen Peng**, Junbin Liang, Sheng Wen, Md Zakirul Alam Bhuiyan, Yiqiao Cai, and Jiannong Cao. “Following Targets for Mobile Tracking in Wireless Sensor Networks.” *ACM Transactions on Sensor Networks (TOSN)*, 12(4): 31:1-31:24, 2016.
- [4] **Zhen Peng**, Tian Wang, Md Zakirul Alam Bhuiyan, Xiaoqiang Wu, and Guojun Wang. “Dependable Cascading Target Tracking in Heterogeneous Mobile Camera Sensor Networks.” *Springer International Publishing, Algorithms and Architectures for Parallel Processing (ICA3PP Workshops and Symposium)*, 2015: 531-540.
- [5] Tian Wang, **Zhen Peng**, Junbin Liang, Yiqiao Cai, Yonghong Chen, Hui Tian, and Bineng Zhong. “Detecting Targets Based on a Realistic Detection and Decision Model in Wireless Sensor Networks.” *Springer International Publishing, Wireless Algorithms, Systems, and Applications (WASA)*, 2015: 836-844.

- [6] Tian Wang, **Zhen Peng**, Yonghong Chen, Yiqiao Cai, and Hui Tian. "Continuous tracking for mobile targets with mobility nodes in WSNs." *International Conference on Smart Computing (SMARTCOMP)*, Hong Kong, pp. 261-268, November 3-5, 2014.

TECHNICAL SKILLS

Programming Languages: C++, C, Python

Frameworks and Tools: AVX-512, OpenMP, MPI, TensorFlow Lite for Micro, MLIR, TVM, TensorRT

AWARDS & HONORS

Student Travel Grands, PACT '18	2018
Student Travel Awards, ASPLOS '18	2018

TEACHING EXPERIENCE

College of William & Mary

Grading TA: CSCI 304 Computer Organization	02/2018 – 05/2018
Grading TA: CSCI 304 Computer Organization	09/2017 – 01/2018
Grading TA: CSCI 243 Discrete Structures	02/2017 – 05/2017
Grading TA: CSCI 141 Computational Problem Solving	09/2016 – 01/2017

PROFESSIONAL SERVICES

Conference Reviewer:

- The 29th International Conference on Information, Communication and Automation Technologies (ICAT-2023)
- The 30th IEEE International Conference on High Performance Computing, Data, and Analytics (HiPC-2023)
- The 2023 International Conference for High Performance Computing, Networking, Storage, and Analysis (SC-2023)
- The 36th IEEE International Parallel & Distributed Processing Symposium (IPDPS-2022)
- The 26th ACM SIGPLAN Symposium on Principles and Practice of Parallel Programming (PPoPP-2021)
- The 35th ACM International Conference on Supercomputing (ICS-2021)
- The 35th IEEE International Parallel & Distributed Processing Symposium (IPDPS-2021)
- The 49th International Conference on Parallel Processing (ICPP-2020)
- The 17th Annual IFIP International Conference on Network and Parallel Computing (NPC-2020)
- 2020 BenchCouncil International Symposium on Benchmarking, Measuring and Optimizing (Bench-2020)
- The 26th IEEE International Conference on High Performance Computing, Data and Analytics (HiPC-2019)
- 2019 BenchCouncil International Symposium on Benchmarking, Measuring and Optimizing (Bench-2019)
- The 16th Annual IFIP International Conference on Network and Parallel Computing (NPC-2019)
- The 5th International Conference on Big Data Computing and Communications (BIGCOM-2019)
- The 28th International Conference on Computer Communications and Networks (ICCCN-2019)
- The 15th Annual IFIP International Conference on Network and Parallel Computing (NPC-2018)
- The 25th IEEE International Conference on High Performance Computing, Data, and Analytics (HiPC-2018)
- The 15th IEEE International Conference on Ubiquitous Intelligence and Computing (UIC-2018)
- The 3th International Symposium on Sensor-Cloud Systems (SCS-2017)